

Computer Systems from the Ground Up



Winter 2025

<https://cs107e.github.io>

Have you ever wondered ...

- how a computer represents data?
- what operations a computer understands?
- how a program executes?
- what happens when a user types on keyboard?
- how text and drawing appears on a display?

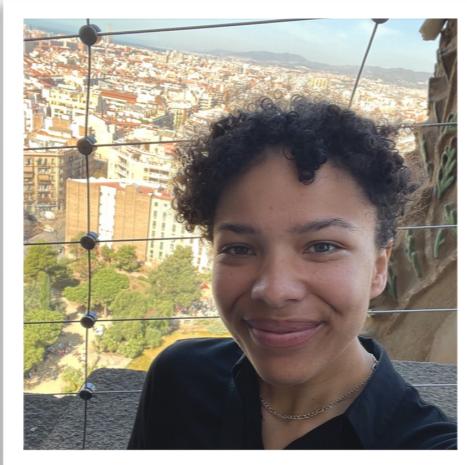
- how things *really* work inside this wondrous box?

These questions and more to be answered by
studying computer systems!

Who?



Ben



Haven



Didi



Julie

+ you!
Intrepid young padawans



Chris

Learning goals

- 1) Understand how computers represent data, execute programs, and control peripherals
- 2) Master your tools

Understanding is empowering!

RISC-V processor and memory architecture

Peripherals: GPIO, timers, UART, ...

Assembly language and machine code

Low-level representation of information / bits

From assembly language to C

Function calls and stack frames

Serial communication and strings

Modules and libraries: Building and linking

Memory management: Memory map & heap

Processor control, interrupts

Master your tools

UNIX command line: bash, cd, ls, ...

Text editor: vim, emacs, sublime, ...

Programming languages: C, assembly

Compiler: gcc

Assembler: as

Linker/loader: ld

binutils: nm, objcopy, objdump, ...

make

git and github.com

documentation: markdown



Syllabus

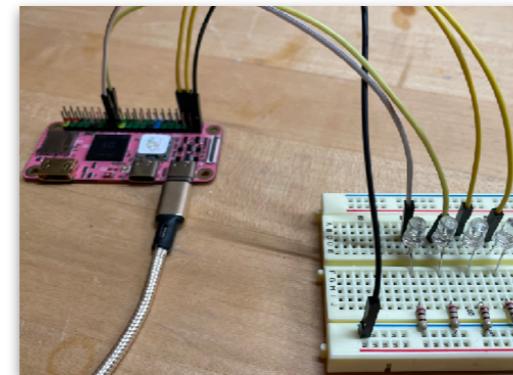
§1 Bare Metal Programming

- RISC-V architecture and assembly language
- C functions and pointers
- Serial communication
- Linking and loading
- Memory allocation



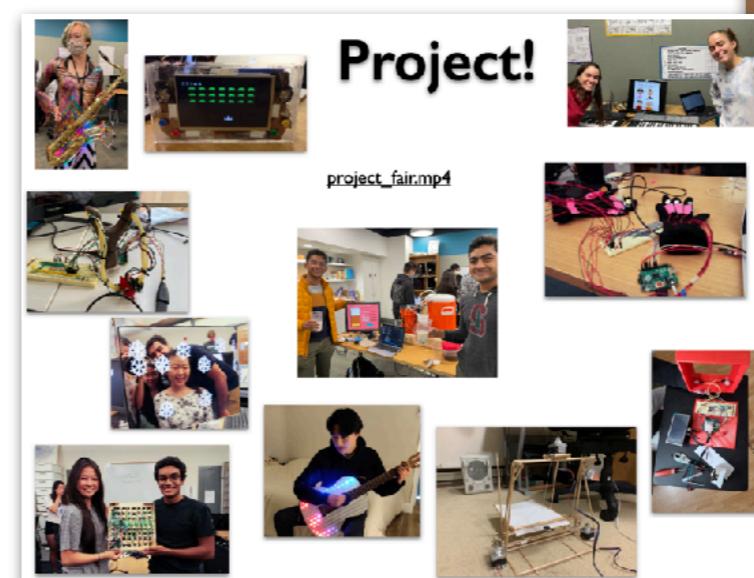
§2 Build a Personal Computer

- Keyboard
- Graphics
- Interrupts



§3 Create Your Own Project

- Sensors
- Performance



Weekly Cadence

Mon	Tue	Wed	Thu	Fri
You are here! → Intro/RISC-V				Assembly
	Lab 0: Setup	Assign 0: Setup		
C Control				C Pointers
	A0 due	Lab 1:ASM	Assign 1:ASM	
				C Functions
MLK Day				
	A1 due	Lab 2: C	Assign 2: Clock	
	A2 due			

<https://cs107e.github.io/schedule/>

Each week has a focus **topic**

Pair of coordinated **lectures** on Fri and Mon

Lab on Tue/Wed evening

Assignment handed out Wed after lab, due following Tuesday 5pm

Staying on pace leads to best outcomes!

Lectures

Attendance is **necessary**

Content is unique to our course, no textbook
The readings/slides not a standalone resource
Lectures not recorded

In-person attendance allows you to participate, ask questions,
keep on schedule, stay connected

Labs

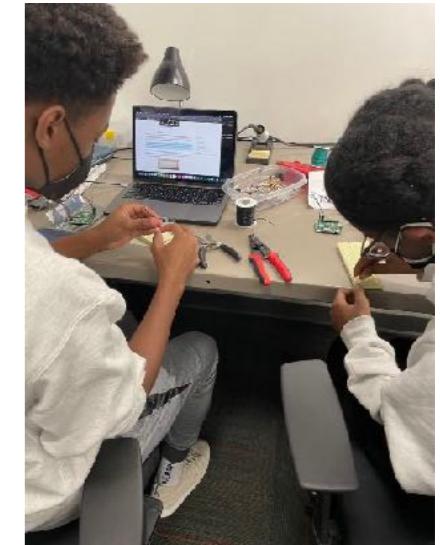
Set of guided exercises, follow up on lecture content

Work in groups

Complete exercises and check in with staff

Leave lab ready to start assignment!

Lab participation is **mandatory**



Philosophy: lab is hands-on, collaborative, supported, **fun!**



Assignments

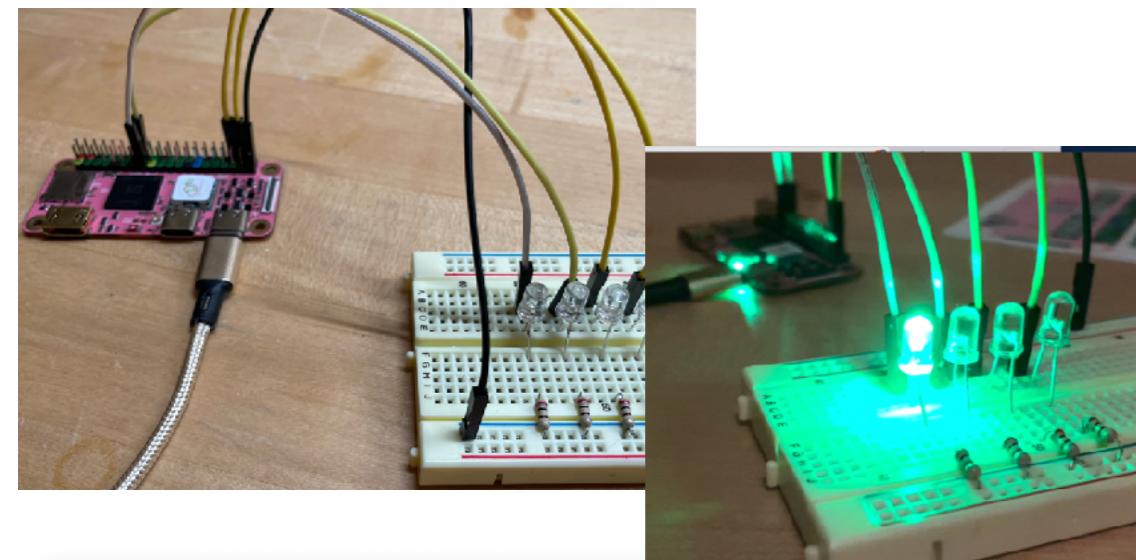
Weekly assignments that build on each other
This is where the learning really happens!

Each assignment has

- **Basic** requirement (tight spec, guided steps)
- Optional **extension** (opportunity for exploration/creativity)

Opportunity to revise/resubmit to correct missed basic requirements

End goal is complete working system of your own



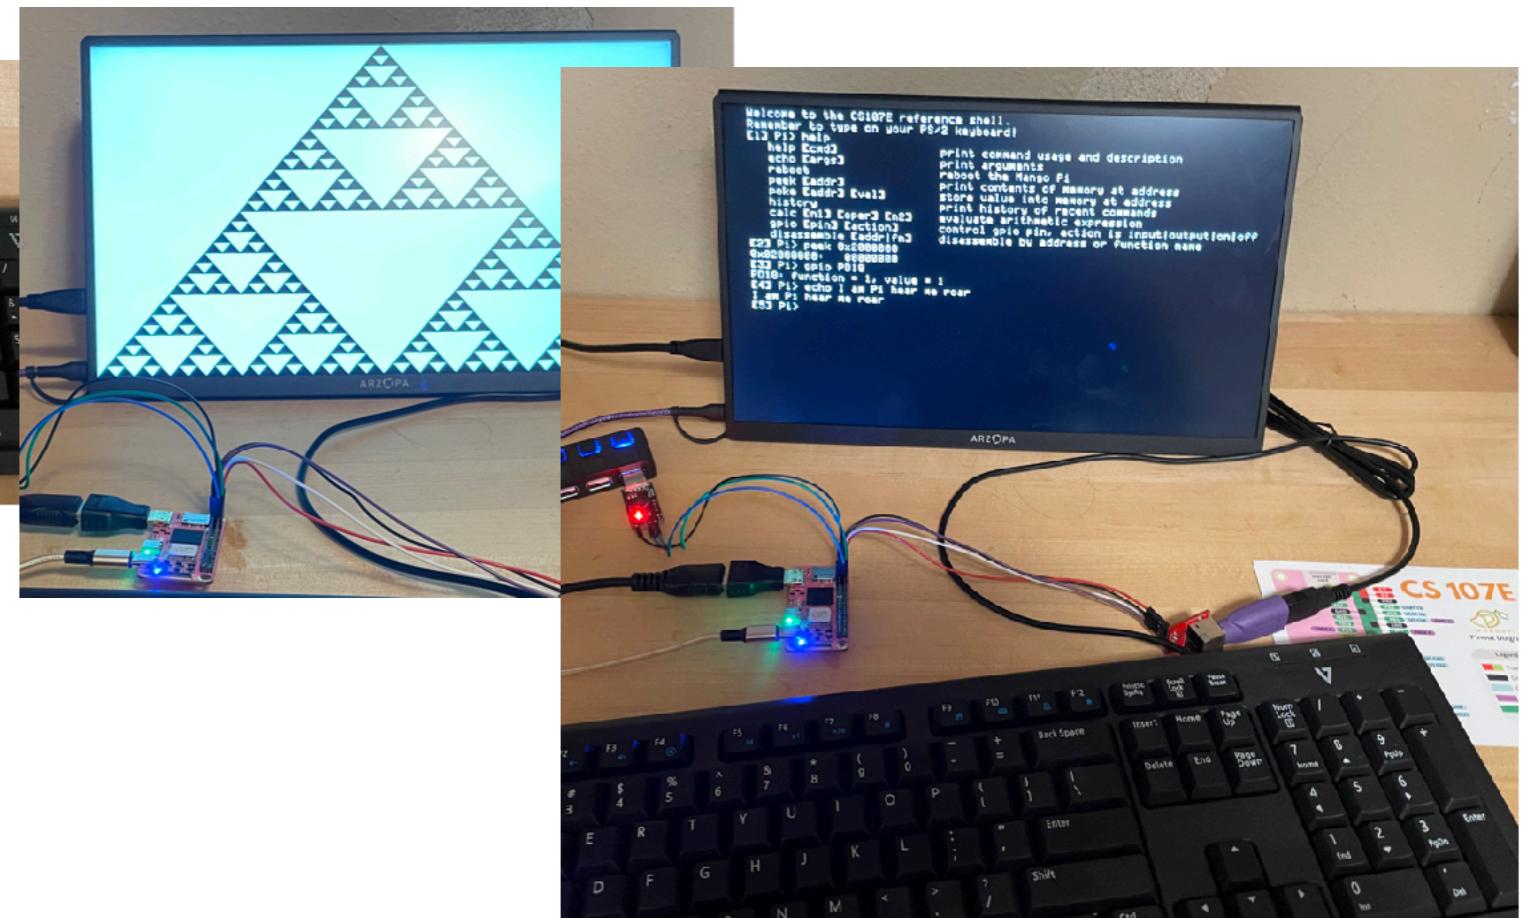
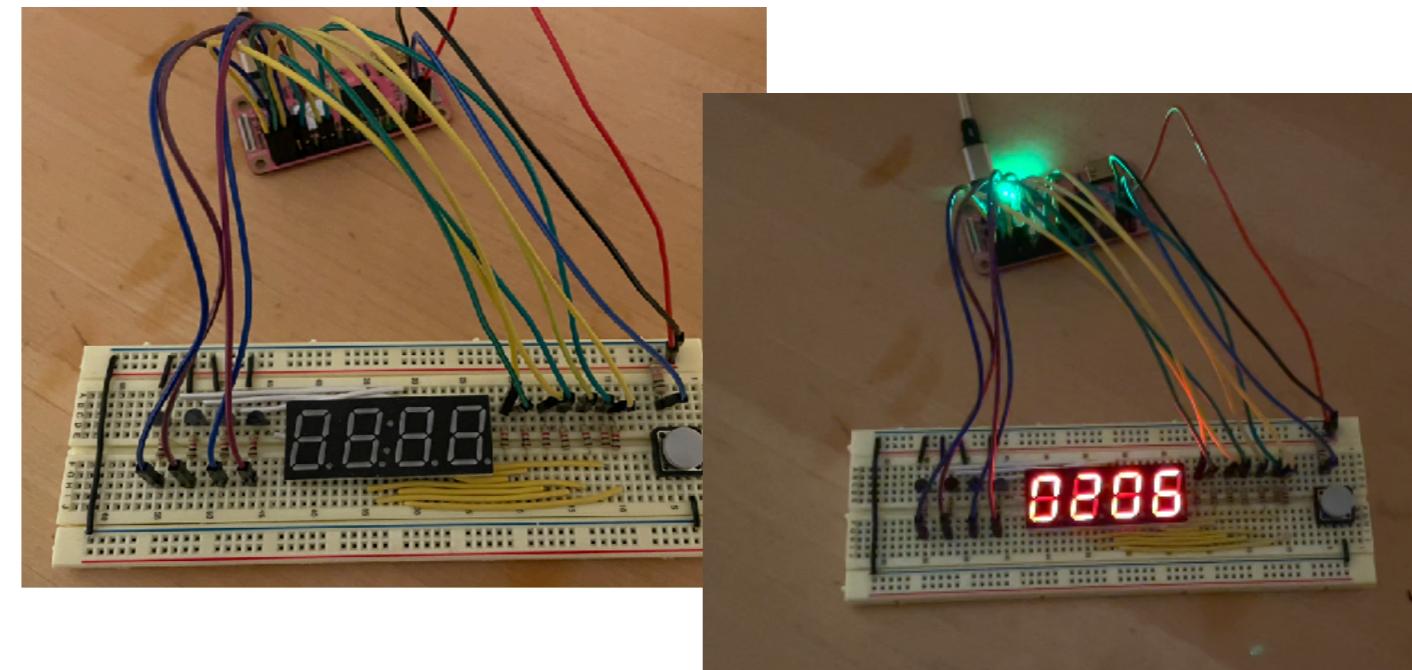
```

Terminal
[OTG] [USB] |-----| HDMI |-----| mini |---0
0---| --| +-----+ |-----+ |-----+
     |   | D1 | micro |
Mango Pi | SoC | | sd | |
     |-----+ +-----+
| @ @ @ @ @ @ @ @ @ @ @ @ @ @ |
| - - - - - - - - - - - - - - |
| @ @ @ @ @ @ @ @ @ @ @ @ @ @ |
0-----0

3V3 01|02 5V
PG13 03|04 5V
PG12 05|06 GND
PB7 07|08 P38 (TX)
GND 09|10 P09 (RX)
PD21 11|12 PB5
PD22 13|14 GND
PB0 15|16 PB1
3V3 17|18 PD14
MOST 19|20 GND
MISO 21|22 PC1
SCLK 23|24 CS0
GND 25|26 PD15
PE17 27|28 PE16
PB10 29|30 GND
PB11 31|32 PC0
PB12 33|34 GND
PB6 35|36 PB2
PD17 37|38 PB3
GND 39|40 PB4

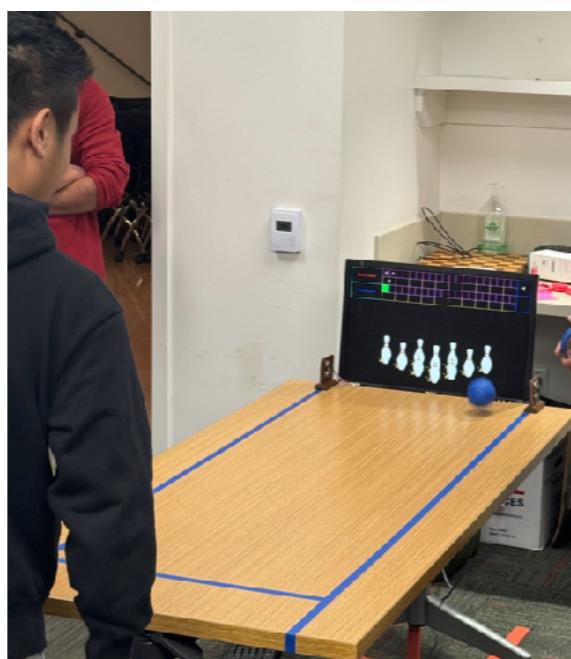
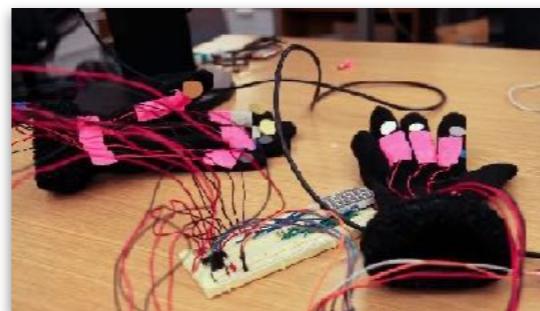
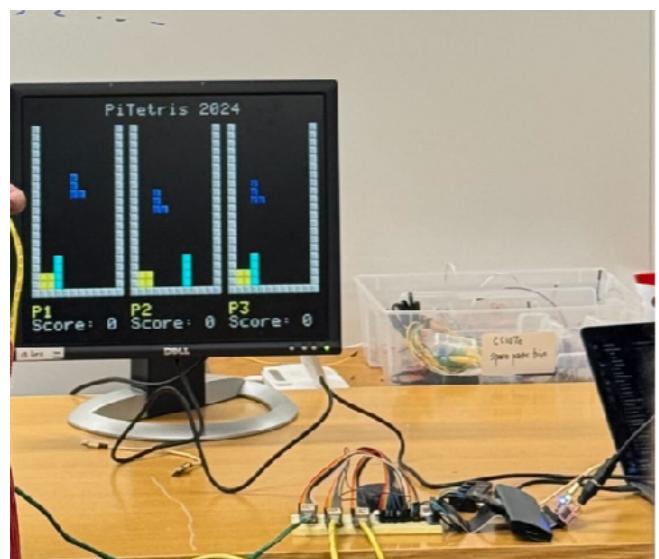
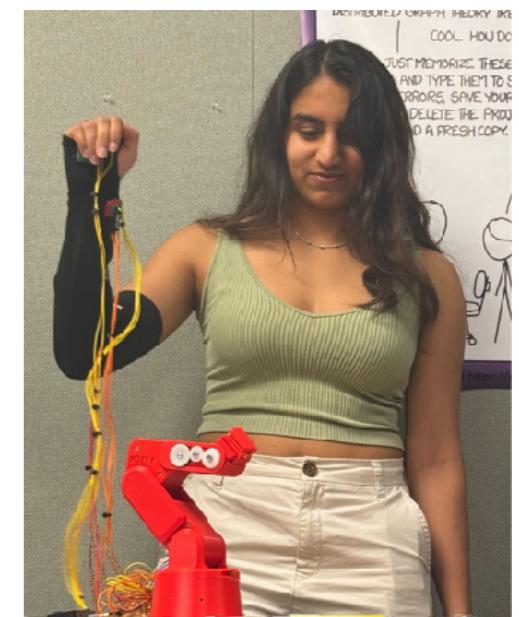
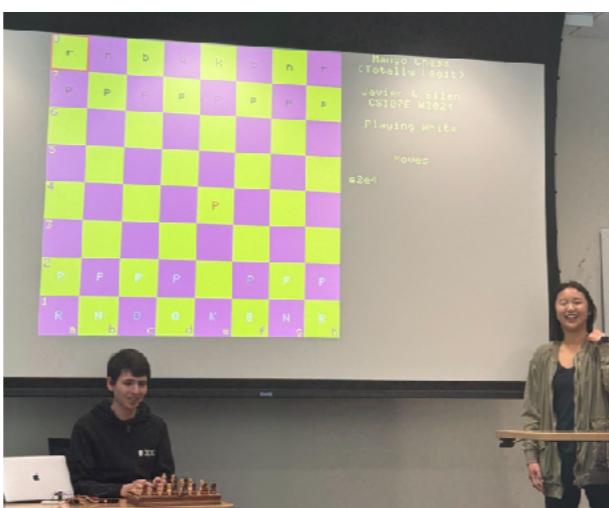
Meta-Z for help | 115200 8N1 | NOR

```



Nearly every instruction is code you wrote yourself!

Projects!



Markers for success

- Necessary prereqs: CSI06B, C++, debugging

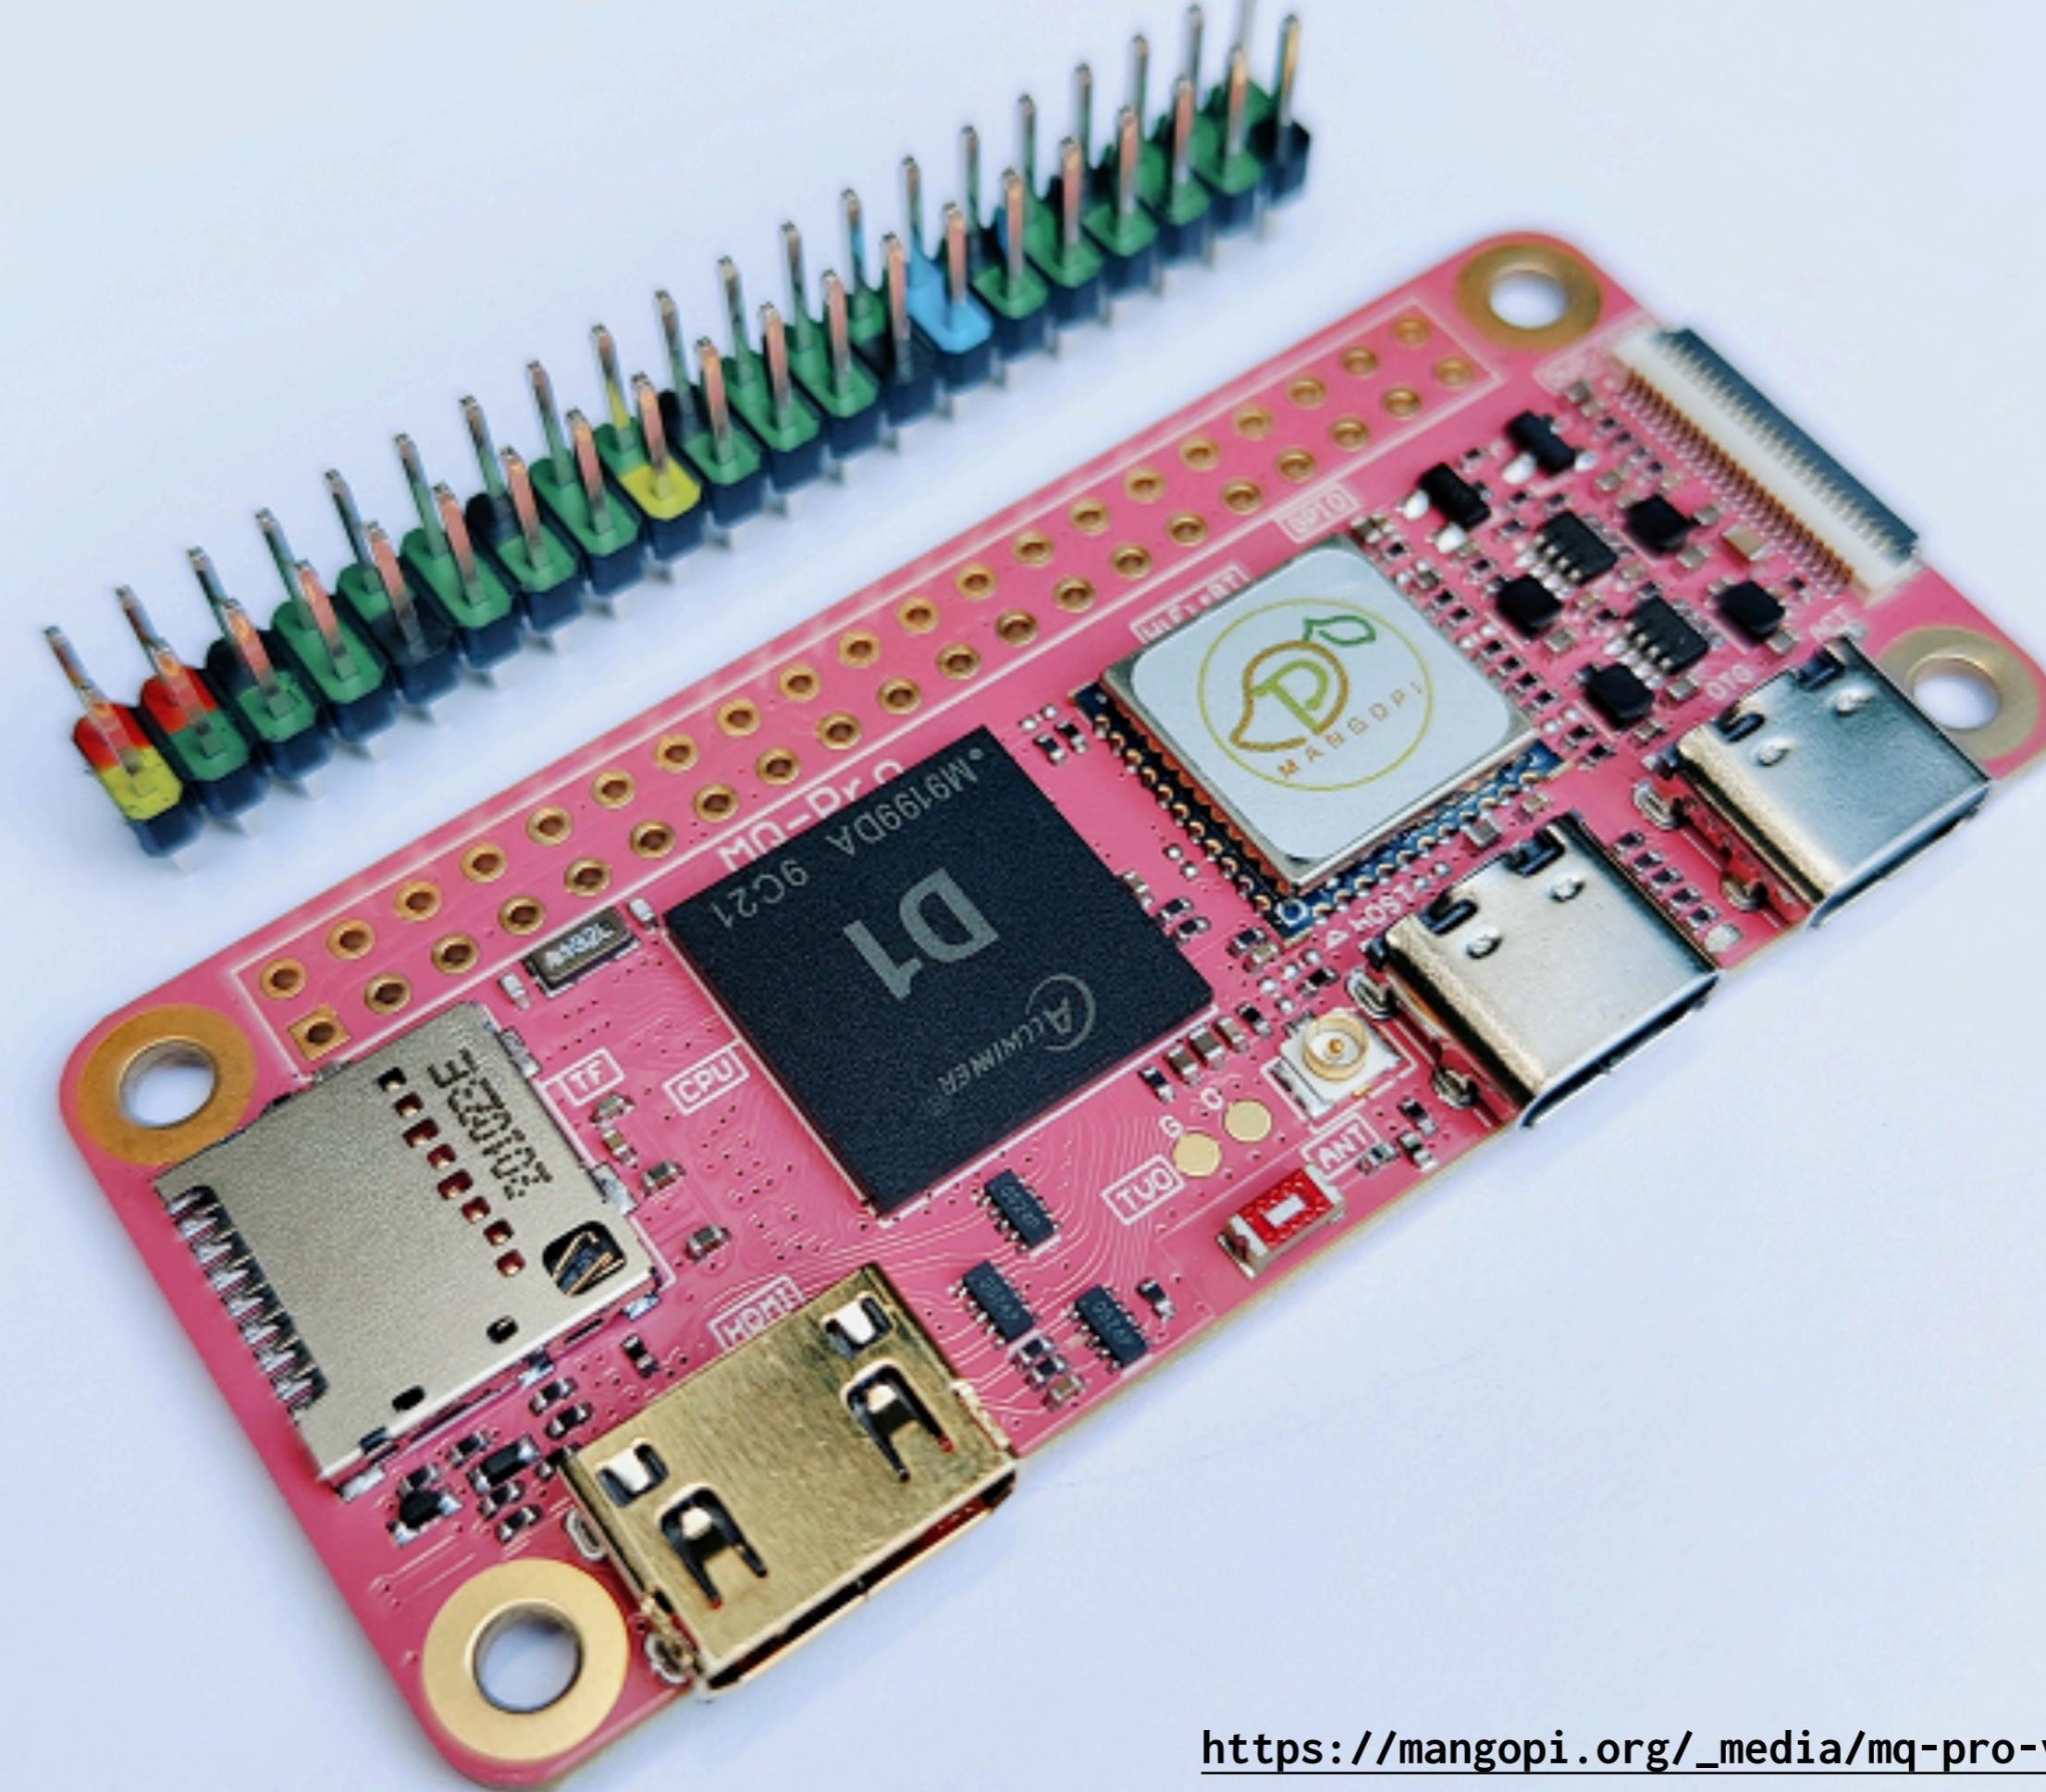
- Curiosity
- Perseverance
- Motivation



Value and appreciate small learning community!

How to thrive in this course

- Consistency, follow through
- Leverage our resources, support, feedback
- Be **curious**. Learn by **doing**. Ask for and offer **help**.



https://mangopi.org/_media/mq-pro-v12-ibom.html

von Neumann architecture

CPU:

ALU

Registers

Control Unit

instruction register, program counter

Memory:

Stores data and instructions

Mechanisms for input/output

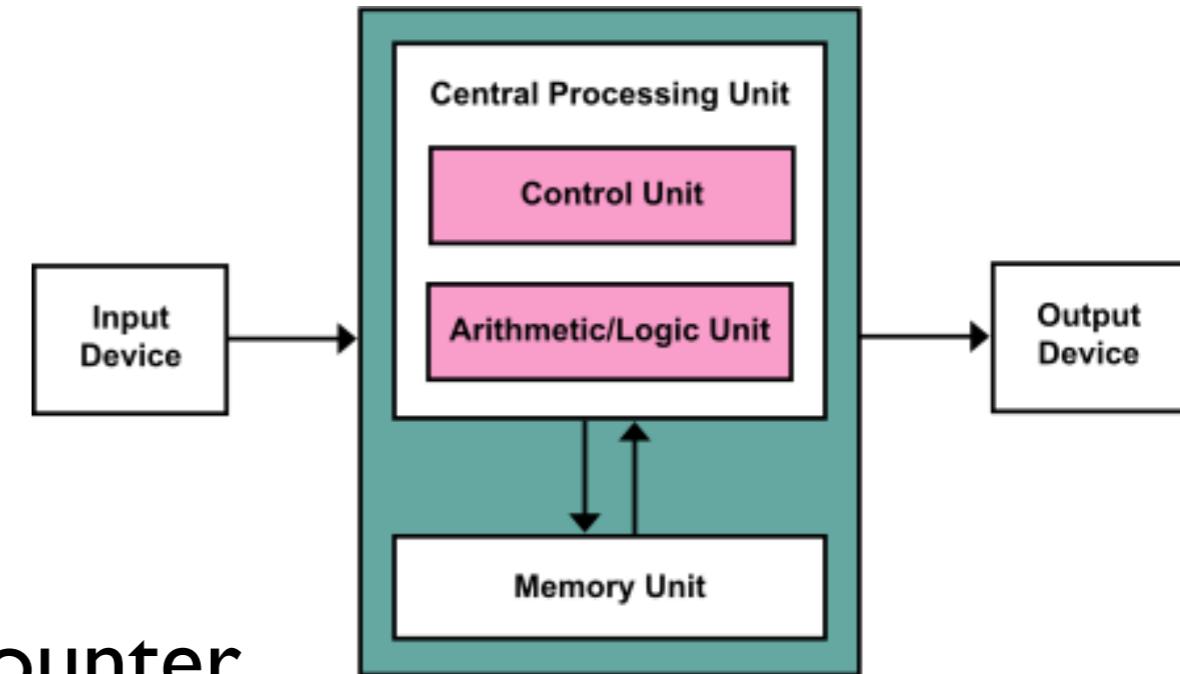


Image credit https://en.wikipedia.org/wiki/Von_Neumann_architecture#/media/File:Von_Neumann_Architecture.svg

Critical innovation:

both instructions and data kept in memory

"stored program" computer

Memory Map

Memory is a large array

Storage locations are accessed
using a 64-bit index, called the
address

Address refers to a *byte* (8-bits)

4 consecutive bytes form a *word*
(32-bits)

Maximum addressable memory is
16EB (But... 42-bit address lines)

1GB physical memory



0xffffffff

$$2^{10} = 1024 = 1 \text{ KB}$$

$$2^{20} = 1 \text{ MB}$$

$$2^{30} = 1 \text{ GB}$$

$$2^{32} = 4 \text{ GB}$$

$$2^{64} = 16 \text{ EB}$$

First week: set up

Before lab:

- Review course guides:
 - Unix tools (shell, editor, git)
 - Electricity (simple circuits, Ohm's law)
 - Number representation (binary, bit operations)
- Install development tools

During lab:

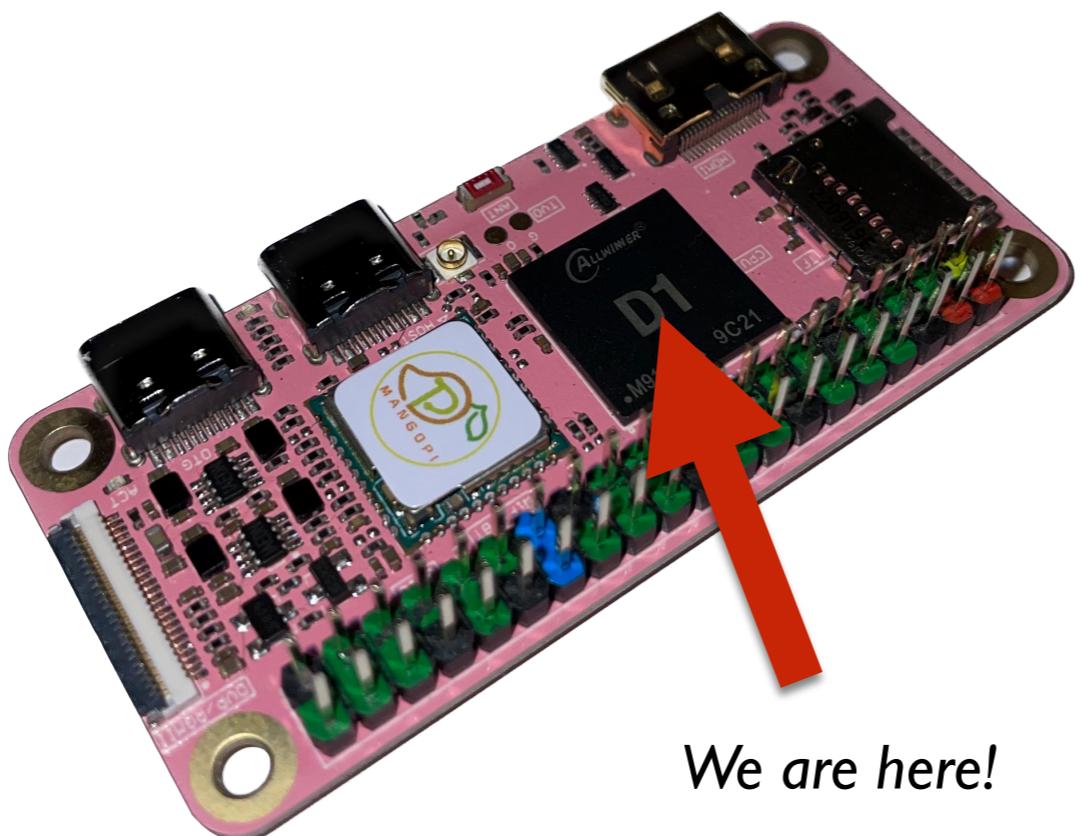
- Establish comfort with background topics
- Practice with environment/tools, build productive habits
- Get help resolving any installation snags
- Meet one another!

Today: RISC-V ISA

Historical milestones in computer architecture

Origins of RISC-V

RISC-V instruction set architecture



We are here!



Whirlwind history tour

Mainframe era (50s,60s)

IBM had 4 incompatible lines of computers

System/360: one ISA to rule them all (backward

compatibility becomes a thing...) now oldest surviving ISA

DEC PDP-8 -> PDP-11 -> Vax (Bell Labs, unix)

ISA prioritize assembly programmer



Personal desktop computer

Intel iAPX 432 6+ year visionary failure (32-bit, OO, Ada)

8086 stopgap developed 1 year (16-bit 8Mhz 29K transistors)

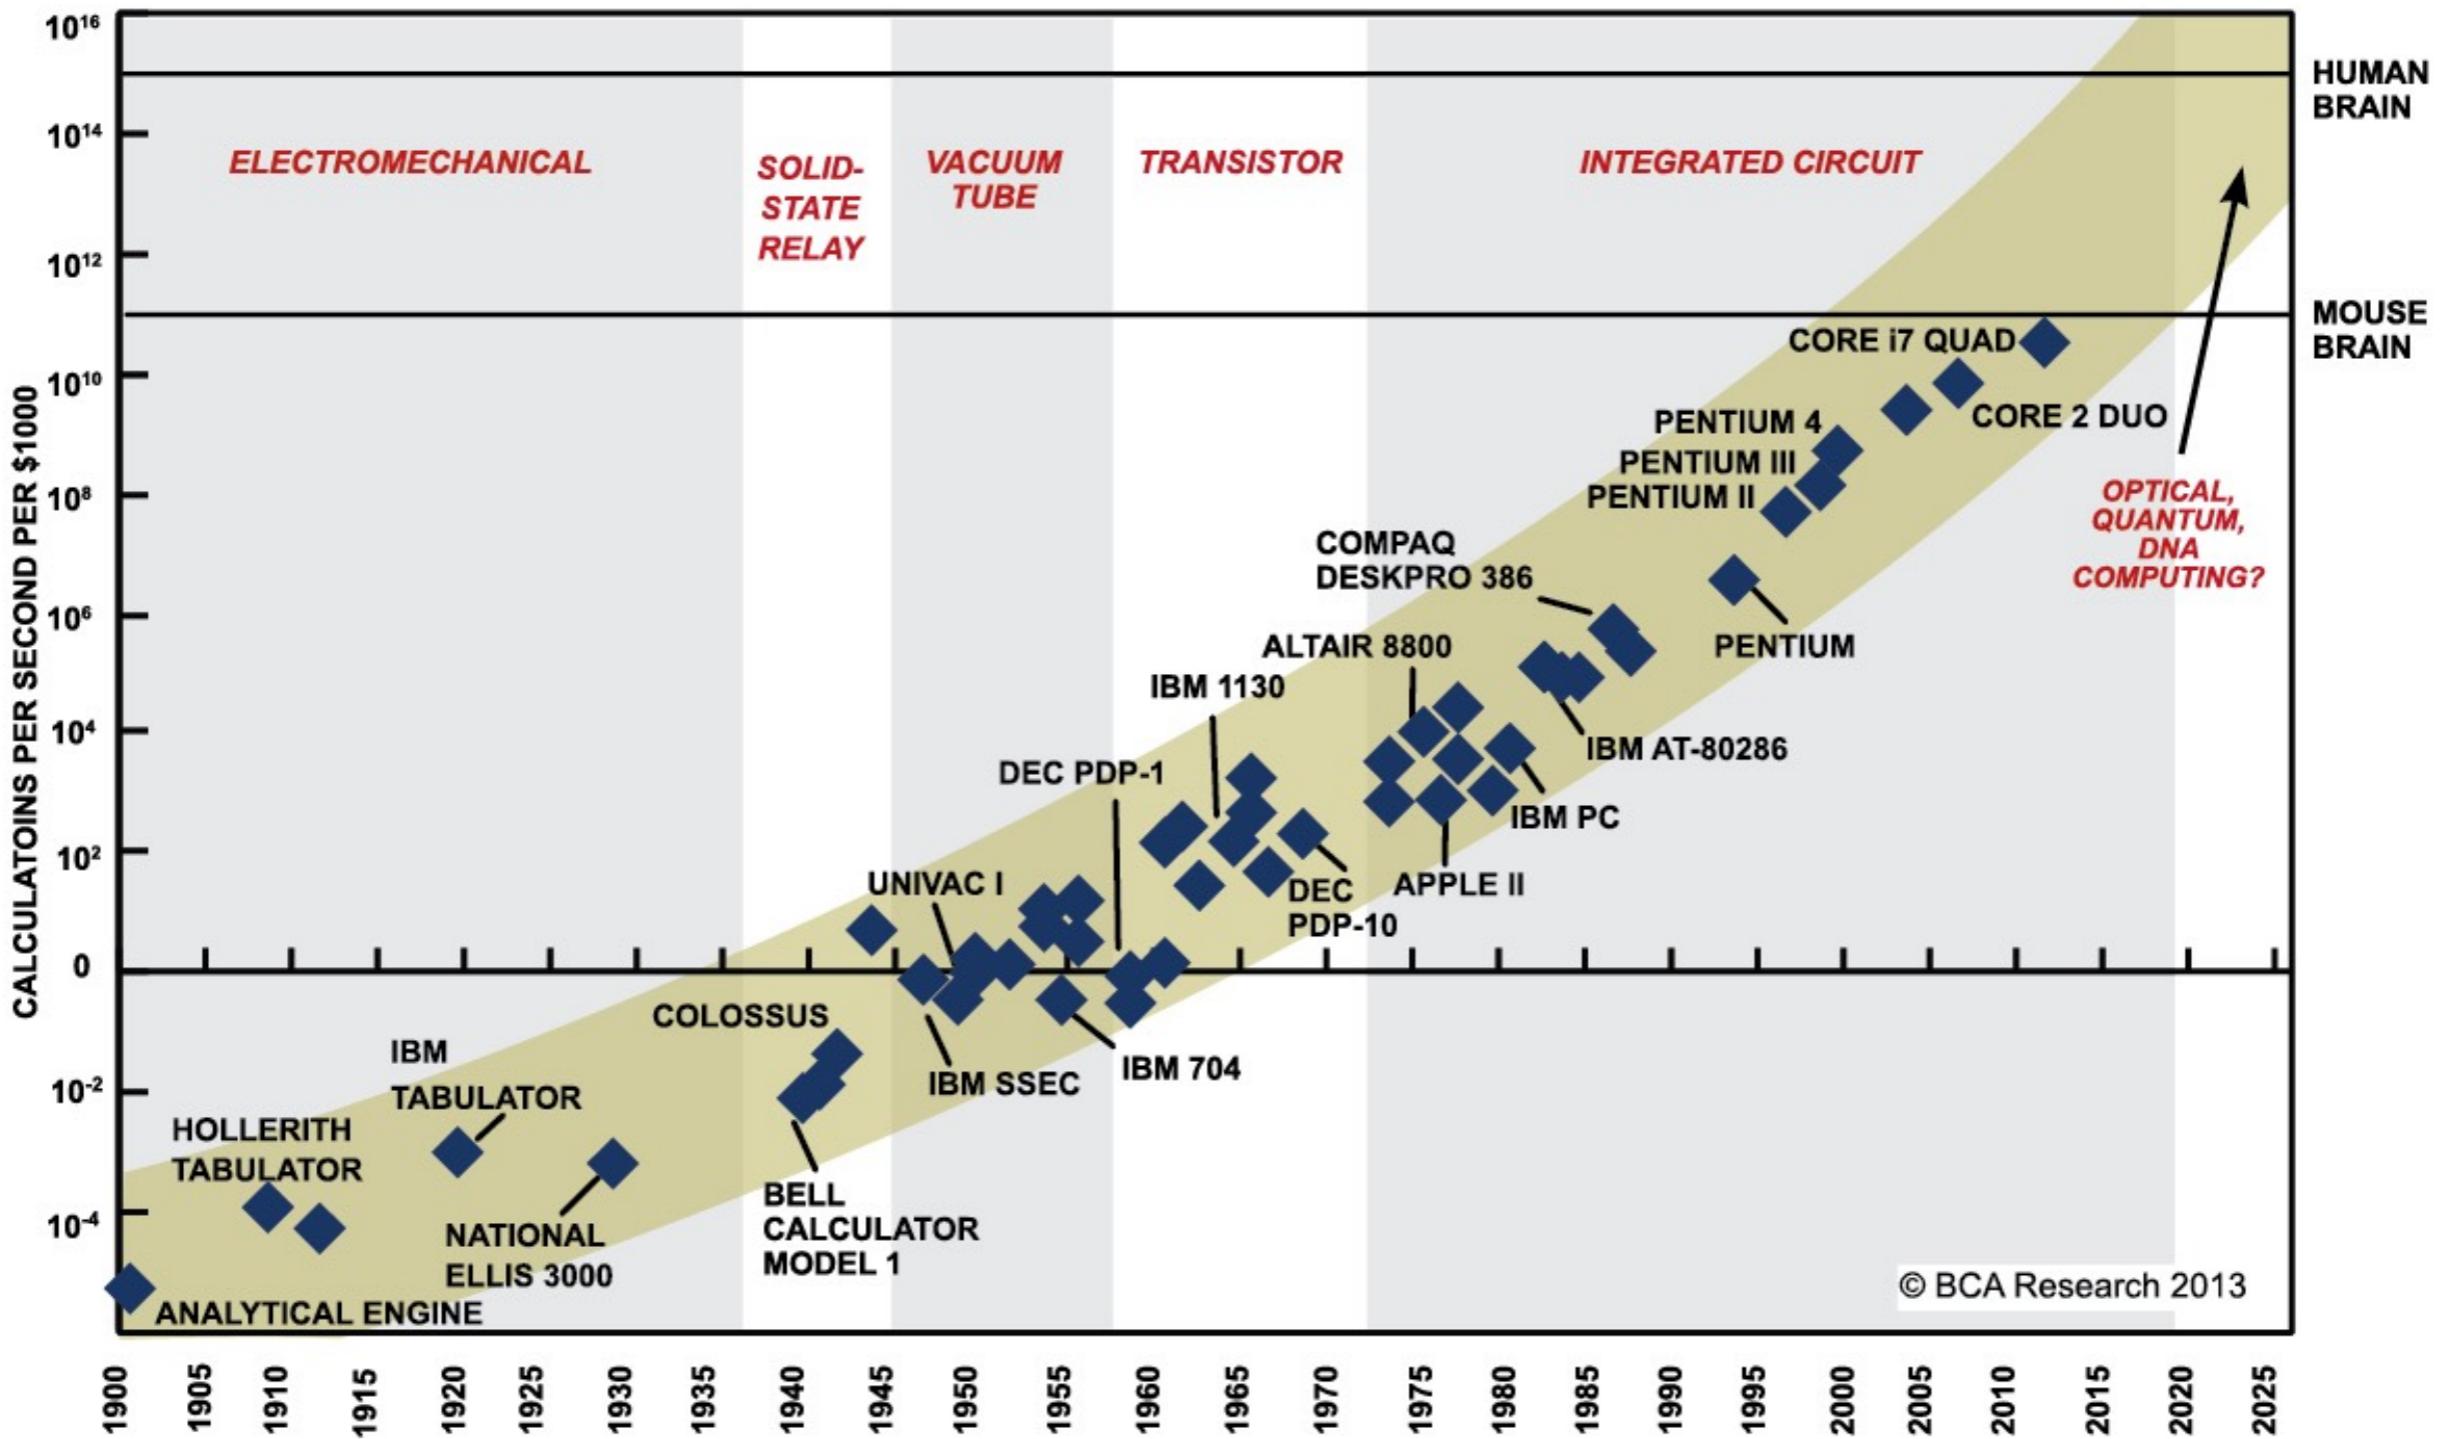
Increasing use of HLL, compiler



Moore's Law

Count of transistors in IC doubles every 2 years

Progression, Moore's law



SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPoints BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.

The “semantic gap”

Gap between HL languages and machine code

One HL construct = many machine instructions

Improve performance of compiled code by adding fancier machine instructions (procedure calls, array access, etc)?

CISC (complex instruction set computer)

Insns to match HL can be complex to implement in hardware

But... researchers in 70s work out:

Compilers mostly don't emit the complex instructions

Avoid special cases, sequence of simple ops often faster anyway

Real-world programs spend most of their time executing simple ops

Complex ops slow down execution, even when you don't use them

From CISC to RISC

Hennessy & Patterson 1982 MIPS

(microprocessor w/o interlocked pipeline stages)

RISC reduce footprint, small number of simple ops

Many advantages to keeping base simple

Simpler to design/verify

Lower costs (die area, higher yield, energy consumption)

Enable pipelined implementation

Higher throughput (faster clock, fewer cycles per ins)

Don't pay for what you don't use

Assembly examples

C code

```
void set(int arr[], int i, int val) {  
    arr[i] = val;  
}
```

x86

0:	8914b7	movl %edx, (%rdi,%rsi,4)
3:	c3	retq

ARM

0:	e7802101	str r2, [r0, r1, lsl #2]
4:	e12fff1e	bx lr

rv

0:	00259593	sll a1,a1,0x2
4:	00b50533	add a0,a0,a1
8:	00c52023	sw a2,0(a0)
c:	00008067	ret

AAA—ASCII Adjust After Addition

Opcode	Instruction	Op/En	64-bit Mode	Compat/Leg Mode	Description
37	AAA	Z0	Invalid	Valid	ASCII adjust AL after addition.

Instruction Operand Encoding

Op/En	Operand 1	Operand 2	Operand 3	Operand 4
Z0	N/A	N/A	N/A	N/A

Description

Adjusts the sum of two unpacked BCD values to create an unpacked BCD result. The AL register is the implied source and destination operand for this instruction. The AAA instruction is only useful when it follows an ADD instruction that adds (binary addition) two unpacked BCD values and stores a byte result in the AL register. The AAA instruction then adjusts the contents of the AL register to contain the correct 1-digit unpacked BCD result.

If the addition produces a decimal carry, the AH register increments by 1, and the CF and AF flags are set. If there was no decimal carry, the CF and AF flags are cleared and the AH register is unchanged. In either case, bits 4 through 7 of the AL register are set to 0.

x86 puts the C in CISC

**arm is RISC
but ...**

*ARM Instructions***A4.1.20 LDM (1)**

LDMIAEQ SP!, {R4-R7, PC}

31	28	27	26	25	24	23	22	21	20	19	16	15	0
cond	1	0	0	P	U	0	W	1		Rn		register_list	

LDM (1) (Load Multiple) loads a non-empty subset, or possibly all, of the general-purpose registers from sequential memory locations. It is useful for block loads, stack operations and procedure exit sequences.

The general-purpose registers loaded can include the PC. If they do, the word loaded for the PC is treated as an address and a branch occurs to that address. In ARMv5 and above, bit[0] of the loaded value determines whether execution continues after this branch in ARM state or in Thumb state, as though a BX (`!loaded_value`) instruction had been executed (but see also *The T and J bits* on page A2-15 for operation on non-T variants of ARMv5). In earlier versions of the architecture, bits[1:0] of the loaded value are ignored and execution continues in ARM state, as though the instruction MOV PC, (`!loaded_value`) had been executed.

Syntax

`LDM{<cond>}<addressing_mode> <Rn>{!}, <registers>`

RISC-V origins

2010 Cal EECS, research iterating on HW design

Existing ISAs no good (complex, proprietary, legacy)

Goal define clean-slate ISA

Substrate for future research

Teaching architecture/systems

Andrew Waterman, Yunsup Lee

Dave Patterson, Krste Asanovic

Design of the RISC-V Instruction Set Architecture

Andrew Waterman



Electrical Engineering and Computer Sciences
University of California at Berkeley

Technical Report No. UCB/EECS-2016-1
<http://www.eecs.berkeley.edu/Pubs/TechRpts/2016/EECS-2016-1.html>

January 3, 2016

<https://people.eecs.berkeley.edu/~krste/papers/EECS-2016-1.pdf>

Design goals

"Universal"

Suitable for microcontrollers up to supercomputer

Data-informed

Benchmarks, measurements, 50 years of experience

Set footprint from intersection not union

Modular, extensible

Small standard base ISA

Defined extensions outside core

Opcode space for custom/specialized extension

Stable

Maintain backward compatibility (extensions frozen)

Evolve via additional extensions

Free & open

Allows collaboration, competition, innovation

Open cores -> secure & trustworthy, design your own

Interface/implementation

ISA

Externally visible aspects

 registers, instructions, data types, control/exceptions

Huge value of standard, no IP hurdles

Microarchitecture

Implementation of ISA

 Logic design, power/performance tradeoffs

Room for innovation, competition

 e.g. Intel & AMD both implement x86, but very different microarchitectures

Abstraction for the win!

How to understand an ISA

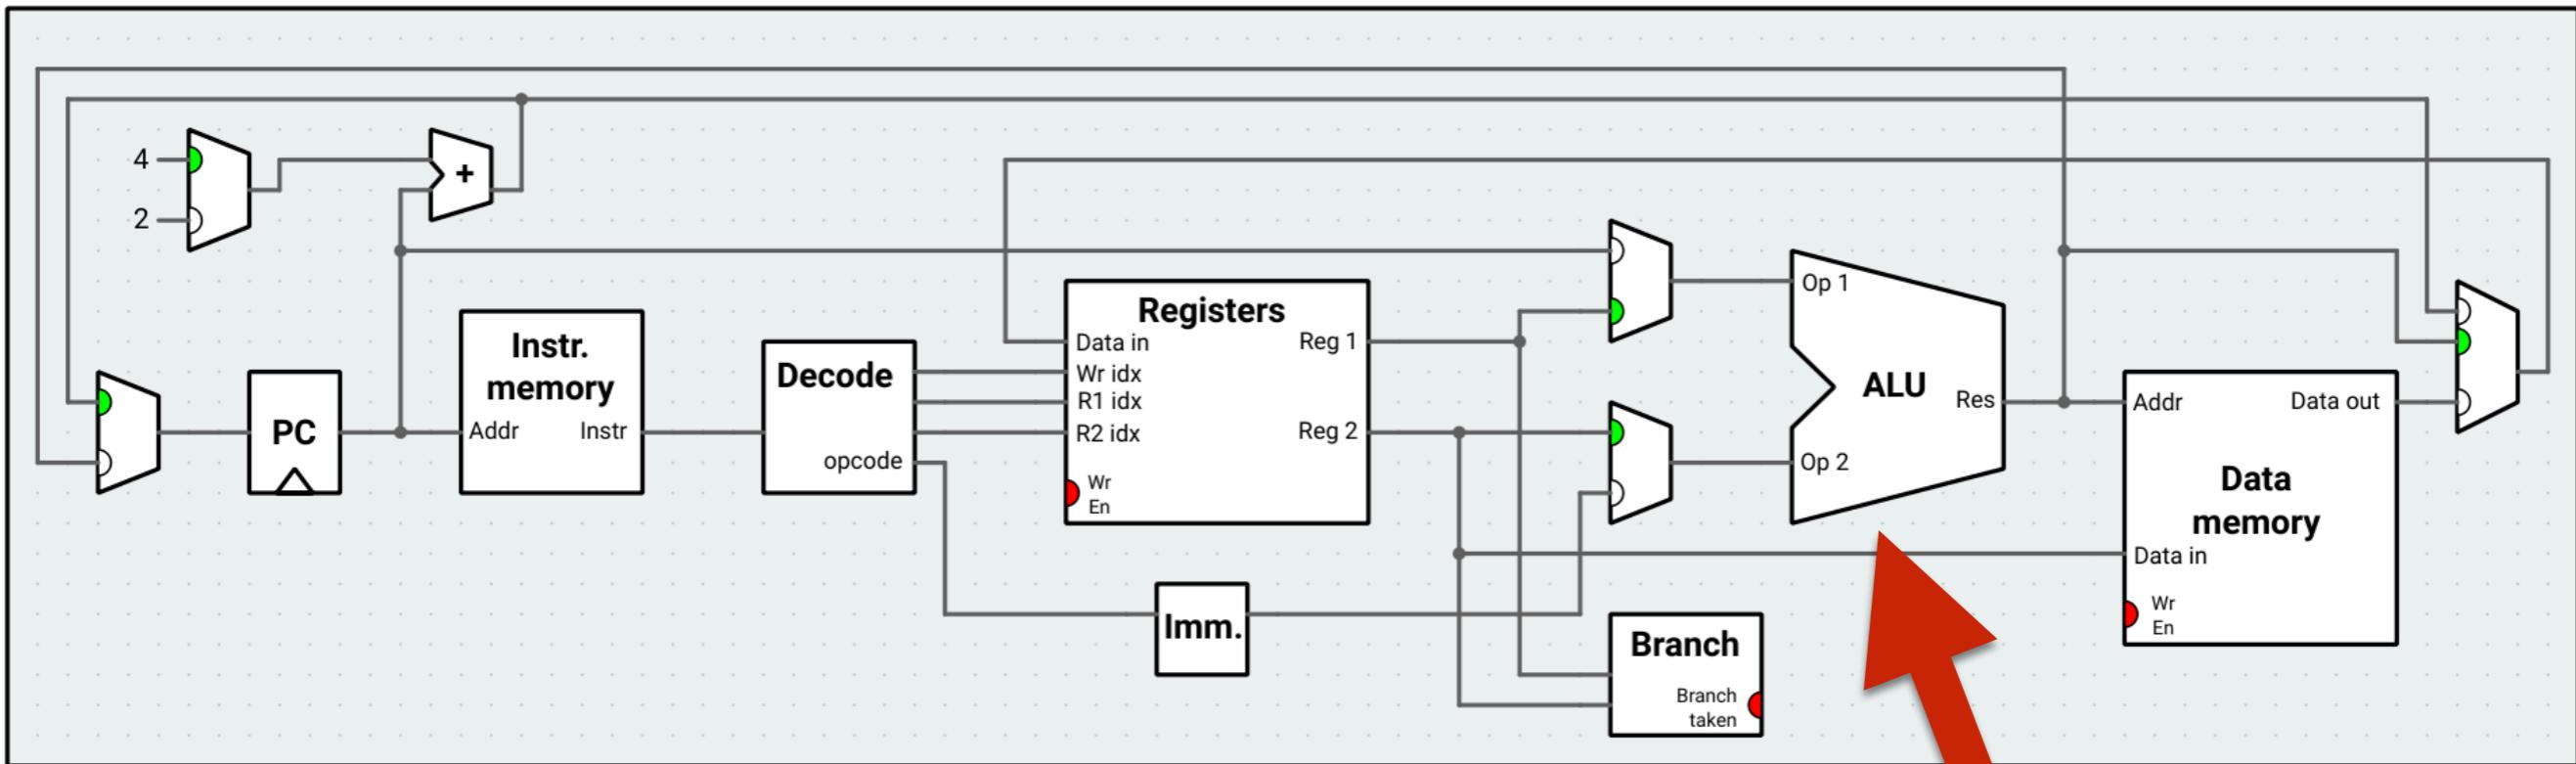
We want to learn how processors represent and execute instructions.

One means of learning an ISA is to follow the data paths in the "floor plan".

Tracing those paths shows where information can flow from/to and how that dictates the operational behavior

Visualizer/simulator can be helpful! <https://ripes.me>

RISC-V Architecture / Floor Plan



<https://ripes.me/>

The "Arithmetic Logic Unit"
This is where arithmetic (e.g. add)
and logic (e.g., and/or) happen

32 registers (x0-x31)

XLEN-1	0
x0 / zero	
x1	
x2	
x3	
x4	
x5	
x6	
x7	
x8	
x9	
x10	
x11	
x12	
x13	
x14	
x15	
x16	
x17	
x18	
x19	
x20	
x21	
x22	
x23	
x24	
x25	
x26	
x27	
x28	
x29	
x30	
x31	
XLEN	

PC = program counter

Example ALU instructions

```
add rd,rs1,rs2  
sub rd,rs1,rs2  
and rd,rs1,rs2  
or  rd,rs1,rs2  
sll rd,rs1,rs2  
srl rd,rs1,rs2
```

R-type (three registers: dest, source1, source2)

```
addi rd,rs,imm12  
andi rd,rs,imm12  
ori  rd,rs,imm12  
slli rd,rs,imm12
```

I-type (source2 is immediate/constant)

Challenge for you all:

Write an assembly program to count the "on" bits in a given numeric value

```
li a0, some-number  
li a1, 0
```

```
// a0 initialized to input value  
// use a1 to store count of "on" bits in value
```



Ripes visual simulator

The screenshot shows the Ripes visual simulator interface. On the left, there's a sidebar with icons for Processor, Cache, Memory, and I/O, and a status bar showing binary values: 1C0, 1010, 01. Below these are tabs for Editor, Processor, Cache, Memory, and I/O.

The main area has tabs for Source code (Assembly selected), Executable code, View mode (Binary and Disassembled selected), and a search/filter icon.

The Source code tab shows the following assembly code:

```
1 li a0, 67
2 li a1, 0
3 loop:
4    andi a2,a0,1
5    add a1,a1,a2
6    srli a0,a0,1
7    bne a0,x0,loop
8
```

The Disassembled tab shows the corresponding machine code and assembly:

Address	Hex	Assembly
0:	04300513	addi x10 x0 67
4:	00000593	addi x11 x0 0
8:	00157613	andi x12 x10 1
c:	00c585b3	add x11 x11 x12
10:	00155513	srli x10 x10 1
14:	fe051ae3	bne x10 x0 -12 <loop>

The Registers panel (gpr) on the right lists the general-purpose registers (x0 to x12) with their aliases and current values:

Name	Alias	Value
x0	zero	0x00000000
x1	ra	0x00000000
x2	sp	0x7fffffff0
x3	gp	0x100000000
x4	tp	0x000000000
x5	t0	0x000000000
x6	t1	0x000000000
x7	t2	0x000000000
x8	s0	0x000000000
x9	s1	0x000000000
x10	a0	0x000000000
x11	a1	0x00000003
x12	a2	0x000000001

Try it yourself! <https://ripes.me>

Instruction encoding

Another way to understand the design of an ISA is to look at how the bits are used in the instruction encoding.

RISC-V uses 32-bit instructions. Packing all functionality into a 32-bits encoding necessitates trade-offs and careful design.

RISC-V instruction encoding

32-bit RISC-V instruction formats

Format	Bit																																						
	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0							
Register/register	funct7							rs2					rs1					funct3			rd				opcode														
Immediate	imm[11:0]												rs1					funct3			rd				opcode														
Store	imm[11:5]							rs2					rs1					funct3			imm[4:0]					opcode													
Branch	[12]	imm[10:5]						rs2					rs1					funct3			imm[4:1]				[11]	opcode													
Upper immediate	imm[31:12]																																						
Jump	[20]	imm[10:1]							[11]	imm[19:12]										rd				opcode															

6 instruction types

Regularity in bit placement to ease decoding

Sparse instruction encoding (room for growth)

ALU encoding

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7		rs2		rs1	funct3		rd		opcode		R-type			
imm[11:0]				rs1	funct3		rd		opcode		I-type			
imm[11:5]		rs2		rs1	funct3	imm[4:0]		opcode		S-type				
imm[12 10:5]		rs2		rs1	funct3	imm[4:1 11]		opcode		B-type				
		imm[31:12]					rd		opcode		U-type			

add x3, x1 , x2

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

000000000000100000010000000010110011
 0 0 2 0 8 B 3

Tip: in python3: print(bin(0x002081b3))

Immediate encoding

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7	rs2			rs1	funct3			rd	opcode				R-type	
imm[11:0]	imm[11:0]			rs1	funct3			rd	opcode				I-type	

imm[11:0]	rs1	000	rd	0010011	ADDI
imm[11:0]	rs1	010	rd	0010011	SLTI
imm[11:0]	rs1	011	rd	0010011	SLTIU
imm[11:0]	rs1	100	rd	0010011	XORI
imm[11:0]	rs1	110	rd	0010011	ORI
imm[11:0]	rs1	111	rd	0010011	ANDI

Your turn!

addi a0, zero, 21

00000001010100000000010100010011
0 1 5 0 0 5 1 3

Key concepts so far

Bits are bits; bitwise operations

Memory addresses (64-bits) index by byte (8-bits), word is 4 bytes

Memory stores both instructions and data

Computers repeatedly fetch, decode, and execute instructions

RISC-V instructions: ALU, load/store, branch

Resources to keep handy

RISC-V one-page guide

Ripes simulator